

Analysing the saleability of cars through Bayesian Networks

Simone Montali - 983806

May 2021

Abstract

Provided with a **dataset** consisting of data about several cars, this research first tries to infer a **Bayesian Network** that accurately describes the problem and its probabilistic background, then, using the generated network, provides an analysis on the **interest** generated by the car for future saleability with respect to different features of the car. The inference is first computed with **variable elimination**, then an approximation through **Rejection Sampling** and **Likelihood Weighting** is compared to the initial results.

Contents

1	Introduction	2
2	Dataset	2
3	Bayesian Network	2
3.1	Structure learning	2
3.2	Parameter learning	3
4	Exact inference	4
4.1	Theoretical background	4
4.2	Variable elimination with pgmpy	4
5	Approximate inference	5
5.1	Theoretical background	5
5.2	Approximate inference in pgmpy	6
6	Conclusion	7

1 Introduction

The car market has been, in the last century, subject to a constant and strong growth. Nowadays, all kinds of features and gadgets have been added to all the car segments, from small city vehicles to luxury ones. Sometimes, though, to get a better grip on what the people prefer, abstracting from buttons, touch screens and custom horn sounds, **a simpler vision is preferable**. This is the reason why the analysis of a rather simple dataset through probabilistic reasoning could actually produce really interesting insights on how cars are chosen by the public.

2 Dataset

This dataset ([click here](#)) was created in 1997 by two researchers of the Joef Stefan Institute in Ljubljana, Slovenia to demonstrate a qualitative multi-criteria decision analysis (MCDA) method for decision making, DEX. The provided data consists of the classification of different features in ≈ 1800 cars. The target attribute is the **car acceptability**, a measure of the interest in the car demonstrated by the public. It is obvious that this dataset is a strong simplification of what actually happens when buying a car, and the fact that it was artificially crafted by someone doesn't help. Nonetheless, **it perfectly fits** the needs of a project like this one. The features that are stated for each car are the following: the buying price, the maintenance price, the number of doors, the seats, the size of the lug boot, the safety. Obviously, the numeric attributes like the prices are **bucketized**. Therefore, the final attribute list is the following:

attribute	values
acceptability	<i>unacc, acc, good, vgood</i>
buying	<i>vhigh, high, med, low</i>
maint	<i>vhigh, high, med, low</i>
doors	<i>2,3,4,5more</i>
people	<i>2,4,more</i>
lug_boot	<i>small, med, big</i>
safety	<i>low, med, high</i>

The data is in CSV format, so a car in the dataset looks like this:

`vhigh,vhigh,3,2,small,high,unacc`

3 Bayesian Network

Bayesian Networks allow us to **economically encode a probability distribution** over a set of variables, stating the **conditional independencies** across the different variables, a rather interesting notion in the analysis of this problem. They are a sufficient and compact specification of the full joint distribution. No Bayesian Network of the problem was available, but luckily the Python library **pgmpy** has exactly what is needed for us to obtain the network and its *Conditional Probability Tables* from the dataset.

3.1 Structure learning

To learn the model structure, i.e. a Directed Acyclic Graph, there are **two different families of techniques** that can be chosen: the constraint-based approaches, which search for a graph structure satisfying the independence assumptions that we observe in the empirical distribution; and the score-based approaches, which define an objective function for different models, and then search for a high-scoring model. [1] The approach I chose is the latter, needing a **scoring function** $s_D : M \rightarrow \mathbb{R}$, which maps a model $x \in M$ to a real score, based on how well it fits to the given dataset. Then, a search strategy traverses the space of all possible models and selects the one with optimal score. Several scoring functions are available, such as *BDeu*, *K2* and *BIC*. The last one, *Bayesian Information Criterion*, is a good asymptotic approximation of the likelihood of the training data. After having defined a scoring function, we can choose different

search strategies, usually *local searches*, i.e. iterative search algorithms that keep improving a solution until a local maxima is reached. The algorithm I chose is **Hill Climb Search**, a rather simple algorithm that is able to explore the search space (while still not being extraordinarily performant) and find a good solution. To learn the network structure through `pgmpy`, we'll first have to import the data in a Pandas DataFrame:

```
import pgmpy
import pandas as pd
cars_data = pd.read_csv('data/car.data', names=["Buying", "Maintenance",
"Doors", "People", "LugBoot", "Safety", "Acceptability"])
```

Then, we'll just use `pgmpy`'s `HillClimbSearch` with a `BicScore`:

```
from pgmpy.estimators import HillClimbSearch
from pgmpy.estimators import BicScore

hc = HillClimbSearch(cars_data, scoring_method=BicScore(cars_data))
best_model = hc.estimate()
```

3.1.1 Learned graph

After some instants, the search finds its best guess of model. We can therefore print its `edges` with `best_model.edges()` obtaining the following result:

```
[('Buying', 'Maintenance'), ('Safety', 'People'), ('Safety', 'LugBoot'),
 ('Acceptability', 'Safety'), ('Acceptability', 'People'), ('Acceptability',
 'Buying'), ('Acceptability', 'Maintenance'), ('Acceptability', 'LugBoot')]
```

which, in a more graphically appealing way, is the following graph:

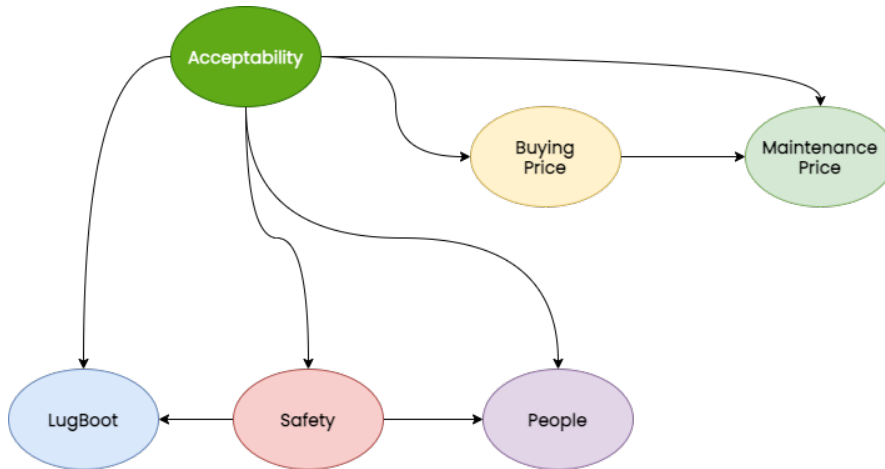


Figure 1: Bayesian network for the cars dataset

3.2 Parameter learning

Now that we've found a good structure for our graph, we'll have to find the values of the Conditional Probability Distributions. The library `pgmpy` is a great tool for this task too. If one would have to find a solution for this problem, it would be pretty natural to think about the **relative frequencies** of the variable states. This approach is, basically, what *Maximum Likelihood Estimation* does: as seen in section 17.1 of [1], the likelihood function for a given choice of parameters θ is the probability the model assigns the training data:

$$L(\theta : \mathcal{D}) = \prod_m P(\xi[m] : \theta) \quad (1)$$

In other words, we want to maximize the probability $P(\text{data}|\text{model})$. This, in practice, is done by computing the state counts and dividing by the conditional sample size. MLE has a high possibility of overfitting on small very datasets. This, fortunately, is not our case, and we would otherwise have needed a different estimator, like Bayesian Parameter Estimation. We can estimate the parameters and print the CPDs with:

```
from pgmpy.models import BayesianModel
from pgmpy.estimators import MaximumLikelihoodEstimator

bayesian_model = BayesianModel(best_model.edges())
bayesian_model.fit(cars_data, estimator=MaximumLikelihoodEstimator)
for cpd in bayesian_model.get_cpds():
    print(cpd)
```

Having finally built our full joint distribution, we can proceed and **answer some questions through inference**.

4 Exact inference

4.1 Theoretical background

4.1.1 Summing out

At this moment, it is very well clear that the basic task for a probabilistic inference system would be, as the word says, performing **inference**, i.e. computing the posterior probability for a set of query variables, given some observed event in the form of evidence variable. Moreover, it is clear that by **summing out** the probabilities from the full joint distribution, we are able to compute any conditional probability. More specifically, a query $\mathbf{P}(X | \mathbf{e})$ can be always computed as:

$$\mathbf{P}(X | \mathbf{e}) = \alpha \mathbf{P}(X, \mathbf{e}) = \alpha \sum_{\mathbf{y}} \mathbf{P}(X, \mathbf{e}, \mathbf{y}) \quad (2)$$

4.1.2 Variable elimination

A crucial improvement in this algorithm is provided by **variable elimination**, which we can simplify as a form of **dynamic programming**, in which the factors get saved into an array to avoid recomputing the same thing multiple times.

4.2 Variable elimination with pgmpy

As always, pgmpy allows us to do so in a straightforward way:

```
from pgmpy.inference import VariableElimination
exact_inference = VariableElimination(bayesian_model)
```

Having instantiated the `exact_inference` object, we can then call queries on it. For example, a nice question to ask would be:

Is maintenance harder for 2-seaters than for family cars?

With just two lines of code, we can produce the probability tables for this query:

```
print(exact_inference.query(["Maintenance"], {'People': "2"}))
print(exact_inference.query(["Maintenance"], {'People': "4"}))
```

Maintenance level	Probability
very high	0.2975
high	0.2595
medium	0.2215
low	0.2215

(a) Maintenance on 2-seaters

Maintenance level	Probability
very high	0.2256
high	0.2450
medium	0.2646
low	0.2648

(b) Maintenance on 4-seaters

It looks like sport cars are not the best kind of choice if you want to stay cheap on oil and filters.

Are cars with high maintenance safe?

Safety	Probability
high	0.2793
medium	0.3240
low	0.3967

It looks like they are not!

5 Approximate inference

Exact inference for **singly connected graphs** (i.e. those in which there's at most one undirected path between any nodes) has a time and space complexity which is linear in the size of the network. This is good for graphs like the famous *Burglary Example*, but once we try inference on **multiply connected graphs**, we lose this nice property. Variable elimination can have **exponential time and space complexity** in the worst case, even when the number of parents per node is bounded. [2] So, what can we do?

5.1 Theoretical background

Approximate inference is the answer. It consists in **randomized sampling algorithms** (also known as *Monte Carlo*) that provide approximate answers having an accuracy depending on the number of samples generated. The obvious, first, step is **generating a sample** from a known probability distribution. This is basically how inference in real life is done: we run the experiment, count the results, compute the probabilities. It is also known as **stochastic approximation**, and for $\hat{N} \rightarrow N$ it will converge to the true probability. The standard deviation of the error in each probability will be proportional to $\frac{1}{\sqrt{n}}$.

5.1.1 Rejection sampling

Rejection sampling is probably the most straightforward way of proceeding. It can be used to compute conditional probabilities $P(X|e)$ by generating samples from the prior distribution, then **rejecting** those that do not match the evidence. Then, pretty easily, the estimate $\hat{P}(X = x|e)$ is obtained by counting how often $X = x$ occurs in the samples.

5.1.2 Likelihood weighting

Rejection sampling has a pretty notable downside: we are generating multiple samples that are inconsistent with our evidence, resulting in a non-trivial amount of wasted time. **Likelihood weighting** tries to solve that. The idea is simple: instead of throwing away the results that don't fit the evidence, we just weigh them less in the final count. The reason for doing that is simple: we would like to sample from the true posterior distribution, but usually this is too hard, as no polynomial-time approximation can exist. Now, a rather big question still has to be

answered: *how do we compute the weights?* The solution is just the product of the conditional probabilities for the evidence variables given their parents:

$$w(\mathbf{z}) = \alpha \prod_{i=1}^m P(e_i \mid \text{parents}(E_i)) \quad (3)$$

This approach can be **much more efficient** than rejection sampling, though it will lose its performance when n becomes higher: the noise that the low-weight samples add becomes more and more *distracting* as the number of samples increases. We know, according to Hoeffding's inequality, that

$$P(|s - p| > \epsilon) \leq 2e^{-2n\epsilon^2} \quad (4)$$

so we may be able to set an accuracy threshold for our approximate inference.

5.2 Approximate inference in pgmpy

As with any other interesting task, `pgmpy` allows us to perform these two. To keep things tidied up, a good idea would be creating a `SampleTester()` class having different methods for RS and LW. The processing method for rejection sampling would look like the following:

```
def process_rs(self, size, evidence, query):
    rejection_sample = self.sampler.rejection_sample(evidence=evidence, size=size,
    return_type='recarray')[query]
    sample_probs = self.return_probs(rejection_sample)
    exact_result = self.exact_inference.query([query], dict(evidence)).values
    absolute_error = self.calculate_error(sample_probs, exact_result)
    return absolute_error
```

while the one for likelihood weighting would be as follows:

```
def process_lws(self, size, evidence, query):
    likelihood_sample = self.sampler.likelihood_weighted_sample(evidence=evidence,
    size=size, return_type='recarray')
    sample_probs = self.return_weighted_probs(likelihood_sample[query],
    likelihood_sample['_weight'])
    exact_result = self.exact_inference.query([query], dict(evidence)).values
    absolute_error = self.calculate_error(sample_probs, exact_result)
    return absolute_error
```

Then, we just have to add some **utilities** for the error calculation and the sample counting:

```
def return_probs(self, samples):
    # Get the unique values and their counts
    unique, counts = np.unique(samples, return_counts=True)
    # Divide the counts by the total number, getting a probability from 0 to 1
    counts = (counts/len(samples))
    # Zip the value and its probability in a dict
    return dict(zip(unique, counts))
def return_weighted_probs(self, samples, weights):
    unique = np.unique(samples)
    # Zero array for the weights sum, which we'll divide by the sum of weights
    counts = np.zeros(len(np.unique(samples)))
    iterator = np.nditer(samples, flags=['f_index'])
    for value in iterator:
        counts[value] += weights[iterator.index]
    counts = (counts/np.sum(weights))
    # Zip the value and its probability in a dict
    return dict(zip(unique, counts))
```

and that's it! We can test our methods with an example evidence *Safety = high, Maintenance = low* and query the acceptability. The **results** are the following:

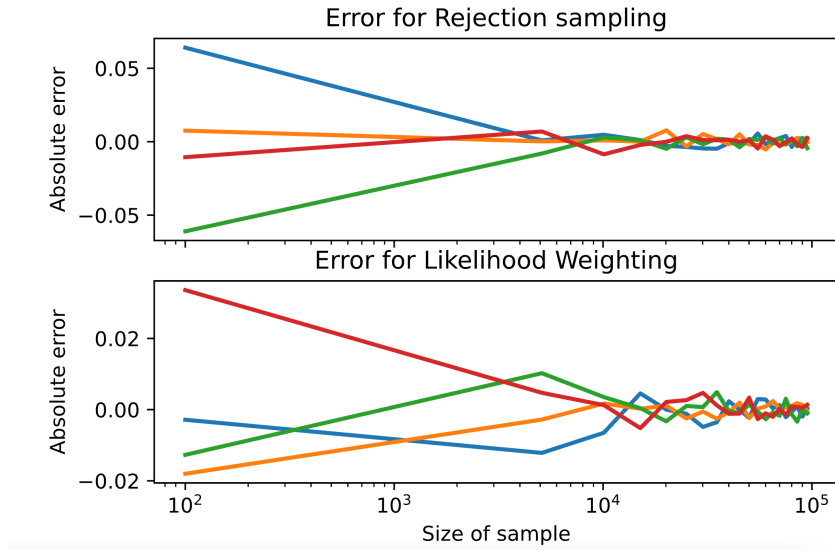


Figure 2: Errors achieved on each acceptability class for the two algorithms.

6 Conclusion

We have seen how Bayesian Networks can prove themselves as an **extraordinary tool** to summarize complex probabilistic domains. This could step up the process of decision making in multiple situations, while still being intrinsically elegant and simple. Tools like `pgmpy` make the whole job a breeze: the programmer just has to think about the problem, the computer will do the rest. That probably is how a truly *intelligent* system would perform. We have then seen how exact inference is - while still being computationally heavy - a great tool, and, moreover, how we can use **approximate inference** to achieve similar results with less, less work.

References

- [1] Nir Friedman Daphne Koller. *Probabilistic Graphical Models: Principles and Techniques*. The MIT Press, 2009.
- [2] Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Prentice Hall, 3 edition, 2010.